Architecting Conversational Shopping Assistants: Integrating Large Language Models with Industrial Product Search Stacks

1. Introduction to the Conversational Commerce Paradigm

The evolution of electronic commerce is currently undergoing its most significant infrastructural shift since the transition from desktop to mobile interfaces. The catalyst for this transformation is the integration of Large Language Models (LLMs) and agentic artificial intelligence into the core product search stack. Historically, digital retail relied on keyword-driven, inverted-index search architectures that placed the cognitive burden of product discovery entirely on the consumer. Users were required to translate their complex, multi-dimensional needs—such as finding a durable, travel-friendly stroller within a specific budget—into truncated keyword strings, subsequently navigating faceted filters and rigid categorization trees to isolate relevant products.

Conversational shopping assistants reverse this dynamic. By deploying autonomous and semi-autonomous AI agents, platforms are shifting from passive retrieval systems to active, goal-oriented orchestrators. These agents do not merely parse queries; they synthesize multi-turn dialogue, track complex state constraints, elicit unstated user preferences, and execute transactions across decentralized architectures. Designing such systems requires coupling the probabilistic, generative nature of large language models with the deterministic, latency-sensitive requirements of an industrial-strength product search stack.

This research report provides an exhaustive analysis of the integration patterns necessary to construct state-of-the-art conversational shopping assistants. The analysis systematically deconstructs the architecture required across query understanding, hybrid retrieval, machine-learning-driven ranking, agentic orchestration, and factual grounding, providing a comprehensive blueprint for the next generation of generative AI retail systems.

2. Landscape Survey of Industry Deployments

The theoretical potential of agentic commerce is currently being validated at an unprecedented scale by major technology and retail conglomerates. Analyzing these vanguard deployments reveals distinct architectural philosophies regarding how generative AI interfaces with existing product graphs and inventory databases.

Platform	Architectural Focus	Key Technical Differentiators
Amazon Rufus	Deep integration of custom LLMs with existing catalog and Web data	Utilizes AWS Trainium/Inferentia chips for parallel decoding to reduce time-to-first-token latency. Heavily relies on retrieval-augmented generation (RAG) and reinforcement learning to synthesize reviews, specs, and Q&A.[1]
Walmart Sparky & Wallaby	Multi-agent orchestration replacing traditional search interfaces	Powered by "Wallaby," a suite of retail-specific LLMs trained on decades of proprietary consumer data. Sparky acts as an end-to-end planner rather than a mere query responder, executing complex tasks like event planning.[5]
Google Shopping AI	Massive-scale query decomposition and synthesis	Integrates Gemini foundation models with the Google Shopping Graph (50+ billion products refreshed 2 billion times per hour). Employs "Query Fan-out" to execute parallel sub-queries across disparate knowledge bases.[9]
Perplexity Shopping	Visual, retrieval-augmented product cards with native checkout	Leverages an exabyte-scale index of 200 billion URLs. Differentiates via seamless transaction capabilities ("Buy with Pro" via PayPal) and rigorous benchmarking

		using its open-source search_evals framework.[12]
Meta AI Shopping	Social-graph-aware contextual recommendations	Embeds conversational product discovery natively within Instagram and WhatsApp. Translates conversational intent into scrollable product carousels, connecting AI-captured buyer intent directly to its Advantage+ advertising infrastructure.[14]

The overarching trend across these deployments is the deprecation of monolithic text search in favor of composable, multi-agent architectures. These systems treat catalog search, dynamic pricing logic, and inventory validation as discrete APIs that the conversational agent dynamically invokes based on the evolving context of the user interaction.

3. Query Understanding (QU) and Contextual Fusion

The foundational layer of any search stack is Query Understanding (QU). Traditional QU pipelines rely heavily on statistical natural language processing (NLP), executing tasks such as spell correction, synonym expansion, and Named Entity Recognition (NER) to map queries to catalog schemas. However, traditional QU is brittle when faced with the syntactic variability, implicit constraints, and multi-turn nature of conversational queries.

3.1 LLM-Driven Intent Classification and Query Fan-Out

In an agentic paradigm, LLMs supersede rigid, rule-based intent classifiers. When a consumer inputs a highly contextual query—for example, "I need a sustainable gift under fifty dollars for a coworker who loves cooking"—the agent must extract not only the explicit entities (price < $50, category: cooking) but also the implicit semantic intent (sustainable, premium feel, gift-appropriate).[17]

The most sophisticated architectural response to this challenge is "Query Fan-out," pioneered within Google's Shopping AI ecosystem.[9] Rather than passing the raw natural language string to a search index, the LLM acts as a deep reasoning layer that decomposes the complex prompt into a multitude of micro-questions.[9] The system then executes these sub-queries in parallel across various retrieval surfaces—including product indices, review databases, and external knowledge graphs—before synthesizing the disparate data streams into a cohesive recommendation.[11] This approach transforms query understanding from a sequential parsing task into a distributed execution plan, drastically increasing the semantic surface area of the search operation.

3.2 Query Rewriting and Expansion Optimization

The application of LLMs for query rewriting bridges the gap between natural language constraints and the rigid syntax required by backend databases. Research indicates that using LLMs to rewrite or expand queries significantly improves the recall of downstream retrieval systems. For instance, the LLM-R2 framework utilizes a large language model to recommend dynamic rewrite rules for database systems, optimizing query execution efficiency without altering the underlying data structure.[20]

Similarly, query expansion leverages the LLM's vast parametric memory to inject highly correlated domain terminology into the search string. Expanding a query from basic keywords to include related concepts (e.g., adding "limited training data" and "domain adaptation" to a query about "low-resourced languages") provides the downstream re-ranker with a broader, yet highly relevant, semantic scope.[22] This technique directly combats the "vocabulary mismatch" problem inherent in traditional inverted indices.

3.3 Conversational Context Fusion and State Tracking

A defining characteristic of a shopping assistant is its ability to maintain context across multiple turns of dialogue. If a user receives a list of laptops and subsequently replies, "show me the silver one, but cheaper," the agent must perform complex coreference resolution. This requires robust Dialogue State Tracking (DST), which monitors the evolving constraints of the conversation.[23]

State-of-the-art implementations utilize models specifically fine-tuned for DST. Recent methodologies involve using frontier models like GPT-4 to simulate thousands of user-agent interactions, which are then used to fine-tune smaller, faster models (e.g., LLaMA 2) for real-time state prediction.[23] These models maintain the conversation state as a continuously updating JSON payload of key-value pairs (slots), tracking user goals, price parameters, and brand exclusions. This state payload is then injected into the structured search_products tool call, ensuring that every subsequent retrieval operation respects the accumulated history of the interaction.

4. Retrieval Infrastructure at Catalog Scale

Because e-commerce catalogs frequently contain hundreds of millions of Stock Keeping Units (SKUs), it is computationally impossible to pass the entire product database into an LLM's context window. Therefore, robust, high-speed retrieval mechanisms are non-negotiable prerequisites for agentic synthesis.[24]

4.1 Modality Evaluation: Sparse, Dense, and Hybrid Retrieval

The industry standard for agentic retrieval is the hybrid model, which fuses the precision of sparse lexical search with the semantic understanding of dense vector retrieval.

Retrieval Pattern	Core Mechanism	Architectural Trade-offs
Sparse (Keyword/BM25)	Inverted index search utilizing term frequency-inverse document frequency.	Unmatched precision for exact part numbers and brand names. However, it is highly brittle to misspellings, synonyms, and conversational paraphrasing.[25]
Dense (ANN Embeddings)	Bi-encoder architecture mapping queries and products to high-dimensional vectors (e.g., via FAISS, ScaNN).	Excellent semantic recall; understands intent even when keywords mismatch. Suffers from the "cold-start" problem for newly added SKUs lacking behavioral interaction data.[26]
Hybrid (Sparse + Dense)	Combines scores from both modalities using Reciprocal Rank Fusion (RRF) or learned interpolation models.	Yields the highest overall relevance. Introduces pipeline complexity, requiring parallel execution and score normalization across differing mathematical scales.[12]
Faceted / Structured Filtering	Hard constraints applied via SQL/NoSQL queries on structured metadata (price, size, stock status).	Absolute necessity in e-commerce to prevent hallucinatory recommendations. Must be decoupled from the semantic scoring pipeline to ensure strict compliance with user constraints.[27]

Implementations of embedding-based retrieval have yielded dramatic improvements in agent accuracy. DoorDash, for example, reported a 90% reduction in LLM hallucinations after transitioning to embedding-based RAG retrieval, proving that providing semantically accurate context is more critical to response fidelity than the raw reasoning capabilities of the generation model itself.[28] DoorDash further enhances this process via Entity Linking, using Approximate Nearest Neighbor (ANN) systems to retrieve the closest taxonomy concepts for a

query before prompting the LLM, thereby reducing noise and preserving the context window.[26]

4.2 Entropy-Based Adaptive Recall Set Sizing

A critical architectural challenge in hybrid retrieval is determining the optimal size of the recall set (K). Over-fetching candidate products saturates the network layer, increases latency, and degrades the signal-to-noise ratio within the LLM's context window. Under-fetching risks missing highly relevant edge cases.

Instacart engineers addressed this by developing an adaptive recall system predicated on query entropy.[27] The system recognizes that broad, informational queries (high entropy) require larger recall sets, while highly specific, transactional queries (low entropy) require tighter constraints. Instacart determines the recall threshold dynamically using the following equation:

$$\text{recall_threshold} = \min\left(M, \max\left(L, M \times \frac{\text{query_product_entropy}}{Q}\right)\right)$$

Where L and M represent the minimum and maximum recall boundaries, and Q serves as the clipping value for entropy.[27] By adaptively calculating this threshold in real-time, the system optimizes the distribution between text and semantic retrieval mechanisms. This intelligent constraint on retrieval volume reduced Instacart's overall search latency by 1.5% while simultaneously improving the mean converting position by 1.7%.[27]

4.3 Data Freshness and Real-Time Ingestion

An AI shopping assistant is fundamentally compromised if it retrieves stale data; hallucinating obsolete prices or recommending out-of-stock inventory irreparably damages consumer trust. Legacy search architectures relying on 24-hour batch processing are inadequate for agentic operations.

To guarantee freshness, platforms are migrating to real-time ingestion frameworks. Pinterest's deployment of a next-generation Change Data Capture (CDC) architecture serves as a prime example. The system separates CDC tables from base tables; utilizing Debezium, Apache Kafka, and Apache Iceberg, the CDC tables act as append-only ledgers that record every database mutation with sub-five-minute latency.[29] Base tables maintain a historical snapshot that is rapidly updated via Spark Merge Into operations. This architecture ensures that when a shopping agent executes a database query to validate stock availability, it accesses a real-time reflection of the inventory state, drastically reducing the probability of transactional failure downstream.[29]

5. Advanced Ranking and Re-Ranking Mechanics

Once the retrieval layer surfaces a candidate set of products, these items must be aggressively re-ranked to maximize precision. In an LLM-augmented search stack, ranking is typically bifurcated into stages: L1 (Initial Recall via BM25/ANN), L2 (Semantic Re-ranking), and L3

(Session-aware Personalization).

5.1 Evaluating Re-Ranking Methodologies

The L2 re-ranking stage is the battleground between traditional machine learning and generative AI. The goal is to accurately score the relevance of document D to query Q.

Approach	Mechanism	Performance and Trade-offs
Pointwise LLM Scoring	The LLM evaluates each query-product pair independently to assign a relevance score.	Highly parallelizable across the candidate set. However, it suffers from severe calibration issues; scores are often noisy, inconsistent, and fail to capture the relative superiority of one item over another.[25]
Pairwise Scoring	The model is prompted to evaluate two items simultaneously and declare a winner.	Strong empirical grounding, but computationally explosive. Evaluating a set of N documents requires $O(N^2)$ comparisons, rendering it economically unviable for real-time, large-scale e-commerce routing.[31]
Listwise LLM Ranking (e.g., RankGPT)	The LLM receives the entire list of retrieved candidates in a single prompt and generates an ordered list.	Achieves state-of-the-art performance (highest NDCG@10) due to its holistic understanding of inter-document relationships.[31] However, it is constrained by context window limits and suffers from high latency and "lost-in-the-middle"

		attention degradation for large candidate sets.[32]
Cross-Encoder Models	Dedicated transformer models that concatenate the query and document, processing them through deep attention layers.	The industry standard for production L2 re-ranking. Models like ZeroEntropy deliver deterministic, well-calibrated scores between 0 and 1. They achieve near-LLM accuracy (e.g., 0.78 NDCG@10) with vastly superior economics (p50 latency of ~130ms, 10-30x cheaper than generative LLMs).[25]

5.2 Mitigating Listwise LLM Latency

While RankGPT (powered by GPT-4) demonstrates superior ability to capture fine-grained semantic ordering, its autoregressive generation latency is a critical bottleneck. E-commerce platforms operate on strict latency budgets, where a 100-millisecond delay can correlate to a 1% drop in gross revenue.[35]

To harness the power of LLMs for ranking without breaching latency service-level agreements (SLAs), researchers have developed two primary architectural optimizations:

1. **Sliding-Window Listwise Re-ranking:** For candidate sets exceeding the optimal context length (e.g., $K = 100$), the system processes overlapping windows of products (e.g., window size 20, stride 10). This prevents context collapse and allows for parallel execution, aggregating the local rankings into a global sort order.[36]
2. **In-Context Re-ranking (ICR):** ICR represents a paradigm shift by abandoning autoregressive text generation entirely. Instead, it leverages the internal attention matrix of the LLM. The system assumes that highly relevant documents receive greater attention weights when the LLM processes the query tokens. Because ICR only requires two $O(1)$ forward passes to extract these attention signals, it cuts latency by more than 60% compared to generative methods like RankGPT, while maintaining competitive NDCG metrics.[33]

The prevailing architectural consensus advocates a tiered strategy: utilize fast, cost-effective cross-encoders to sort the top 100 candidates from the retrieval layer, and reserve expensive, listwise LLM evaluation strictly for the top 5 to 10 items immediately prior to final response synthesis.[25]

6. Agent Architecture and Multi-Agent Orchestration

A critical realization in the development of conversational shopping assistants is that a single, monolithic "do-everything" agent inevitably fails.[37] A generalized agent tasked with understanding intent, invoking pricing APIs, checking inventory, and enforcing safety guidelines rapidly exhausts its context window and reasoning capacity, leading to hallucination and workflow collapse.[38] Consequently, the industry is pivoting toward multi-agent orchestration.

6.1 Orchestration Frameworks

Different shopping intents require vastly different execution topologies. The orchestration layer dictates how an agent reasons and navigates the search stack.

- **Single-Turn RAG:** Optimal for basic informational lookups (e.g., "What is the warranty on this television?"). The query is embedded, relevant text chunks are retrieved, and the LLM synthesizes a direct answer. It lacks the capacity for multi-step problem solving.
- **ReAct (Reason + Act):** The standard paradigm for conversational agents. The LLM interleaves internal "thought" traces with external action executions. It observes the output of a tool call (like search_products), reasons about the result, and decides whether to formulate a new query or respond to the user.[39] While highly adaptable, ReAct agents often get trapped in local minima during complex tasks, repeatedly executing failed queries without broader strategic correction.[40]
- **Plan-and-Execute:** The system splits the cognitive load. A "Planner" agent analyzes the user's prompt to generate a comprehensive, step-by-step execution roadmap. An "Executor" agent then processes these steps. This is highly efficient for complex, deterministic workflows, such as cross-referencing multi-item compatibility.[41]
- **Language Agent Tree Search (LATS):** A cutting-edge framework designed for high-stakes, high-complexity reasoning. LATS unifies reasoning, acting, and planning by integrating Monte Carlo Tree Search (MCTS) into the agent's trajectory.[39] Tested extensively on the "WebShop" environment—a simulated e-commerce platform containing 1.18 million real-world Amazon products—LATS allows the agent to explore multiple decision branches simultaneously. It evaluates intermediate rewards using semantic feedback, enabling it to learn from experience and backtrack from unpromising search paths. In benchmark testing, LATS significantly outperformed traditional ReAct prompting in both overall score and task success rate, making it the premier architecture for sophisticated shopping assistants.[39]

6.2 Tool Definitions and Function Calling Schemas

To orchestrate these workflows, the agent must interact with backend APIs via strictly typed structured outputs (typically JSON).[44] A robust shopping assistant requires a minimum baseline of explicitly defined tools:

- search_products(query: str, filters: dict, top_k: int): Executes the hybrid retrieval pipeline.

- get_product_details(product_id: str): Fetches granular Product Detail Page (PDP) attributes, avoiding the injection of full catalog data into the prompt.[44]
- compare_products(product_ids: list, attributes: list): Triggers specialized synthesis engines to evaluate differences between specified variants.
- check_inventory_and_pricing(product_id: str, location_id: str): Queries real-time CDC databases to validate availability before final recommendation.

6.3 Agentic Commerce Protocols: MCP, UCP, and ACP

Hardcoding function calls is difficult to scale across millions of disparate merchants. To solve the interoperability challenge, the industry is rapidly coalescing around three primary agentic communication protocols, each governing a different layer of the commerce stack.[45]

Protocol	Developer / Backers	Domain and Capabilities	Architectural Mechanics
Model Context Protocol (MCP)	Anthropic, AI Developers	**Data Discovery:** Connects AI agents to external tools, catalogs, and APIs.	A client-server architecture where the LLM (MCP Host) dynamically queries an MCP Server to discover available tools and resources (e.g., product catalogs, review databases) at runtime. It focuses strictly on data access and context maintenance, excluding transaction execution.[46]
Agentic Commerce Protocol (ACP)	OpenAI, Stripe	**Decision & Initiation:** Agent-mediated purchase initiation.	Focuses on delegated intelligence. The AI agent helps the user make a buying decision and initiates the order

			on their behalf. It features delegated payment authorization bounds to ensure the AI operates within user-defined financial constraints.[45]
Universal Commerce Protocol (UCP)	Google, Shopify, Walmart	**Transaction Standardization:** Checkout execution and merchant interoperability.	An open-source, decentralized standard that establishes a common language for carts, inventory, and checkout. It allows agents to securely execute transactions across any UCP-compliant merchant without navigating bespoke website DOMs. The retailer remains the merchant of record.[46]

A fully matured shopping assistant utilizes a synthesis of these protocols: it leverages **MCP** to securely ingest real-time product schemas and specifications into its reasoning engine; it utilizes **ACP** rules to navigate the user's purchasing intent and budget; and finally, it executes the transaction seamlessly via **UCP** pipes, utilizing the Agent Payments Protocol (AP2) to ensure cryptographic security during financial settlement.[45]

7. Factual Grounding, Hallucination Prevention, and UX Synthesis

The primary barrier to consumer adoption of LLM shopping assistants is the erosion of trust caused by hallucinations. Agents that fabricate specifications, invent non-existent promotional pricing, or assure users of false compatibility metrics introduce massive liability.

7.1 Guardrail Architectures

Preventing hallucinations requires implementing deterministic constraints atop the probabilistic LLM.

- **GraphRAG for Entity Relationships:** Traditional vector search retrieves text based on semantic similarity, which can strip away hierarchical context. GraphRAG blends vector retrieval with knowledge graph traversals.[55] By explicitly mapping product taxonomies (e.g., "Motherboard X requires Socket Y"), the system prevents the LLM from synthesizing physically impossible or incompatible product pairings.[55]
- **Semantic Tool Selection:** To reduce error rates and token costs, systems utilize vector-based filtering to dynamically restrict the tools available to the LLM based on the user's intent.[55] If a user asks "what are the dimensions of this sofa?", the orchestrator removes the execute_checkout or modify_order tools from the context window, completely eliminating the risk of the agent initiating an unintended transaction.[55]
- **Automated Reasoning Checks:** Advanced enterprise deployments, such as Amazon Bedrock, deploy formal verification techniques. Rather than relying on probabilistic scoring to detect errors, Automated Reasoning Checks utilize mathematical logic to validate the LLM's generated content against the provided domain knowledge, delivering up to 99% accuracy in detecting ungrounded claims.[57]

7.2 Citation Architecture: The "Cite Before You Speak" Paradigm

Trust requires attribution. Amazon researchers recently published a framework establishing a "citation experience" within conversational shopping agents. The core hypothesis is that users must be able to visually verify the origin of an LLM's claim by clicking a citation mark that links directly to the source text (e.g., a specific user review or technical manual).[58]

Deploying citations at scale introduces severe UX and latency challenges. Amazon solved this by implementing In-Context Learning (ICL) coupled with a Multi-UX-Inference (MUI) system. The MUI layer acts as an orchestration mechanism that appends source citations to the LLM output in real-time, without interfering with the rendering of other UX elements (like product carousels or comparison tables).[58] Rigorous online A/B testing demonstrated that enforcing this citation generation paradigm substantially improved the factual grounding of the LLM's responses by 13.83%, while increasing overall customer engagement by 3% to 10%.[60]

7.3 Synthesizing Structured Comparison Tables

A frequent transactional intent involves product comparison. Generating an accurate, side-by-side comparison table from unstructured product descriptions and thousands of divergent user reviews requires advanced synthesis techniques.

This is achieved through Topic-Oriented Multi-Document Summarization (TOMDS).[62] The architecture executes a two-stage process:

1. **Extractive Stage:** The system scans the retrieved product documents and reviews, isolating paragraphs relevant to specific topics (e.g., "battery life," "build quality") and

sorting them by relevance.[62]

2. **Abstractive Stage:** A transformer-based decoder utilizes an external discourse parsing module to map both micro-level (within-paragraph) and macro-level (inter-paragraph) semantic relationships. Guided by a topic-aware attention mechanism, the LLM contrasts the identified attributes across the different products.[62] The output is a highly structured, hallucination-free comparison table that accurately distills complex qualitative differences into a scannable format, dramatically reducing the consumer's cognitive load.[62]

8. Evaluation Frameworks and System Metrics

Deploying these complex architectures requires rigorous, specialized benchmarking frameworks that measure both retrieval fidelity and generative quality.

8.1 Retrieval and Ranking Metrics

Standard information retrieval metrics—such as NDCG@K (Normalized Discounted Cumulative Gain) and Recall@K—remain foundational. NDCG is particularly vital as it evaluates the system's ability to rank the most highly relevant items at the absolute top of the list, heavily discounting relevant items that appear lower in the UI.[64]

However, generic retrieval datasets fail to capture the nuances of commerce. The industry standard benchmark is the **Amazon ESCI dataset** (Exact, Substitute, Complement, Irrelevant).[65] Unlike binary relevance scoring (relevant vs. non-relevant), ESCI acknowledges the reality of shopping: if a user searches for a specific Sony headphone model that is out of stock, a highly rated Bose headphone (a *Substitute*) is a vastly superior retrieval result than a Sony charging cable (a *Complement*), even though the cable shares the exact brand keyword.[65] Evaluating LLM query rewriting and re-ranking models against the ESCI multi-class framework ensures that the AI optimizes for actual purchasing behavior rather than mere lexical overlap.

8.2 Factual Accuracy and Grounding Evaluation

To systematically evaluate whether an agent is hallucinating, organizations are adopting benchmarks like the DeepMind FACTS (Factual Accuracy and Citation Tracking System) framework.[67] FACTS evaluates the precision, relevance, and verifiability of the citations generated by the model. It calculates an F1 score that balances the presence of correct citations (precision) against the coverage of relevant factual statements supported by citations (recall).[69] High F1 scores on the FACTS benchmark confirm that the LLM is tightly adhering to the RAG context payload.

8.3 The Latency Budget and Infrastructure Economics

The most prohibitive barrier to scaling conversational agents is the compute cost and execution latency. A standard retrieval pipeline incurs severe latency cascades: embedding the query (20ms), vector search for candidates (200-300ms), filtering by date/stock (50ms),

cross-referencing (100ms), and cross-encoder re-ranking (150ms). This results in nearly a full second of latency before the LLM generates a single token of text.[70]

Adding standard LLM generation time pushes the total response time beyond acceptable thresholds, directly eroding revenue. Therefore, hyper-scalers rely heavily on hardware acceleration and algorithmic optimization. Amazon Rufus, for example, utilizes AWS Trainium chips paired with vLLM to manage highly efficient, multi-node inference clusters.[3] By applying hybrid parallelism strategies (tensor and data parallelism), the model is sharded across multiple accelerators to maximize memory bandwidth.[3]

Furthermore, to compress text generation times, systems implement **draft-centric speculative decoding** (parallel decoding). A smaller, highly efficient draft model predicts the next sequence of tokens, which the massive foundation model then validates in parallel.[4] This technique practically doubled the inference speed of Amazon's Query Planner models, ensuring that the conversational interface remained responsive even during the massive traffic spikes of Prime Day.[4]

9. Key Technical Challenges and Future Research Milestones

While the foundational architecture of agentic commerce is solidifying, several critical challenges continue to drive industry research.

9.1 Overcoming the Cold-Start Personalization Problem

With the deprecation of third-party cookies and an increase in guest checkouts, agents frequently encounter "cold-start" users who possess no historical behavioral data.[71] Traditional collaborative filtering fails under these conditions.

Conversational agents circumvent this via active preference elicitation. The Conversational Sales (CSales) framework treats the interaction as a strategic dialogue, dynamically inferring contextual user profiles by asking targeted, multi-attribute clarifying questions.[73] Memory architectures utilizing frameworks like LangChain paired with Vector Databases (e.g., Pinecone) allow the agent to extract these elicited preferences, embed them, and persist them across long-term memory structures.[75] When the user returns, the agent retrieves these embedded preferences to instantly tailor the search parameters, effectively eliminating the cold-start deficit.[75]

9.2 Research Milestones and Trajectory

To push the boundaries of current implementations, future research must focus on the following milestones outlined in the architectural blueprint:

1. **Literature and Baseline Synthesis:** Continuous benchmarking of LLM-based query rewriting against traditional, highly optimized Elasticsearch/Solr QU pipelines.

2. **Protocol Interoperability:** Prototyping extensive Tool-call RAG agents that seamlessly span MCP for data discovery and UCP for transaction execution without dropping state logic.
3. **Latency vs. Relevance Economics:** Deep empirical testing of In-Context Re-ranking (ICR) and distilled listwise models (e.g., LiT5-Distill) against traditional LTR to find the mathematical optimum between NDCG gains and compute cost per query.
4. **Multi-turn Constraint Adherence:** Longitudinal evaluation of how effectively an agent maintains strict negative constraints (e.g., "no leather," "under $100") across a 20-turn conversational session utilizing simulated human-agent datasets like PSCon.[76]

10. Conclusion

The integration of Large Language Models into the e-commerce search stack represents a fundamental redesign of the retail architecture. The industry has unequivocally moved past simplistic, single-turn RAG chatbots. Modern conversational shopping assistants—exemplified by deployments like Amazon Rufus, Walmart Sparky, and Google Shopping AI—are sophisticated, multi-agent orchestration engines.

Building these systems requires abandoning monolithic text search in favor of distributed execution. Robust architectures must couple the deep semantic reasoning of LLMs with high-speed hybrid retrieval, entropy-adaptive recall, and fast cross-encoder re-ranking to satisfy strict latency budgets. Furthermore, to operate autonomously across the open web, these agents must adopt standardized communication frameworks such as the Universal Commerce Protocol (UCP) and the Model Context Protocol (MCP). Finally, consumer trust must be guaranteed by enforcing strict factual grounding via neurosymbolic guardrails and explicit, verifiable citation mechanics. As these architectural patterns mature, agentic commerce will cease to be a novel interface and will establish itself as the primary infrastructure governing global digital retail.

Works cited

1. The New Amazon Rufus - Andrew Bell, accessed April 8, 2026, https://44462615.fs1.hubspotusercontent-na1.net/hubfs/44462615/The%20New%20Amazon%20Rufus%2C%20by%20Andrew%20Bell%20&%20Azoma.pdf
2. The technology behind Amazon's GenAI-powered shopping assistant, Rufus, accessed April 8, 2026, https://www.amazon.science/blog/the-technology-behind-amazons-genai-powered-shopping-assistant-rufus
3. How Amazon scaled Rufus by building multi-node inference using AWS Trainium chips and vLLM | Artificial Intelligence, accessed April 8, 2026, https://aws.amazon.com/blogs/machine-learning/how-amazon-scaled-rufus-by-building-multi-node-inference-using-aws-trainium-chips-and-vllm/
4. How Rufus doubled their inference speed and handled Prime Day traffic with AWS AI chips and parallel decoding | Artificial Intelligence, accessed April 8, 2026,

https://aws.amazon.com/blogs/machine-learning/how-rufus-doubled-their-inference-speed-and-handled-prime-day-traffic-with-aws-ai-chips-and-parallel-decoding/
5. Walmart's Sparky: Agentic AI and the Future of Shopping | Kellogg School of Management, accessed April 8, 2026, https://www.kellogg.northwestern.edu/academics-research/research/detail/2026/walmarts-sparky-agentic-ai-and-the-future-of-shopping/
6. What Walmart's AI Architecture Reveals About the Future of Work | by Vijayan Seenisamy, accessed April 8, 2026, https://vijayan-seenisamy.medium.com/what-walmarts-ai-architecture-reveals-about-the-future-of-work-701a3640578e
7. Walmart's Wallaby LLM - A Guide to what it is and how Walmart Sellers can optimize for it - Ecomtent: AI, accessed April 8, 2026, https://www.ecomtent.ai/blog-page/walmarts-wallaby-llm---what-it-is-and-how-walmart-sellers-can-optimize-for-it
8. Walmart Reveals Plan for Scaling Artificial Intelligence, Generative AI, Augmented Reality and Immersive Commerce Experiences, accessed April 8, 2026, https://corporate.walmart.com/news/2024/10/09/walmart-reveals-plan-for-scaling-artificial-intelligence-generative-ai-augmented-reality-and-immersive-commerce-experiences
9. Inside Google Shopping AI mode: How it works and what powers it - Productsup, accessed April 8, 2026, https://www.productsup.com/blog/inside-google-shopping-ai-mode-how-it-works-and-what-powers-it/
10. AI transforms shopping in Search - Think with Google, accessed April 8, 2026, https://business.google.com/us/think/search-and-video/google-shopping-ai-mode-virtual-try-on-update/
11. What is Query Fan-Out in SEO (AI SEO) - Hobo, accessed April 8, 2026, https://www.hobo-web.co.uk/what-is-query-fan-out-in-seo-aiseo/
12. Architecting and Evaluating an AI-First Search API, accessed April 8, 2026, https://research.perplexity.ai/articles/architecting-and-evaluating-an-ai-first-search-api
13. Shop like a Pro - Perplexity, accessed April 8, 2026, https://www.perplexity.ai/hub/blog/shop-like-a-pro
14. Meta Enters the AI Shopping Wars to Challenge ChatGPT and Gemini - CNET, accessed April 8, 2026, https://www.cnet.com/tech/services-and-software/meta-new-ai-shopping-tool/
15. Meta AI Shopping Tool: How It Works & What Brands Must Know (2026) - ALM Corp, accessed April 8, 2026, https://almcorp.com/blog/meta-ai-shopping-tool/
16. Meta's AI Shopping Tool: Conversational Commerce Insights, accessed April 8, 2026, https://i10x.ai/news/meta-ai-shopping-tool-analysis
17. How agentic commerce is becoming the new front door to retail - Microsoft, accessed April 8, 2026, https://www.microsoft.com/en-us/industry/blog/retail/2026/02/09/how-agentic-commerce-is-becoming-the-new-front-door-to-retail/

18. Mastering Query Fan-Out: Optimizing Content for the AI-First Search Era | by RohitM, accessed April 8, 2026, https://medium.com/@RohitM4/mastering-query-fan-out-optimizing-content-for-the-ai-first-search-era-76ae48de3760
19. What Is Query Fan-Out? How One Query Becomes 12 in AI Search | Ekamoira Blog, accessed April 8, 2026, https://www.ekamoira.com/blog/query-fan-out-original-research-on-how-ai-search-multiplies-every-query-and-why-most-brands-are-invisible
20. LLM-R2: A Large Language Model Enhanced Rule-based Rewrite System for Boosting !ery E"iciency - VLDB Endowment, accessed April 8, 2026, https://www.vldb.org/pvldb/vol18/p53-yuan.pdf
21. LLM-R2: A Large Language Model Enhanced Rule-based Rewrite System for Boosting Query Efficiency - arXiv, accessed April 8, 2026, https://arxiv.org/html/2404.12872v1
22. How to Enhance RAG Pipelines with Reasoning Using NVIDIA Llama Nemotron Models, accessed April 8, 2026, https://developer.nvidia.com/blog/how-to-enhance-rag-pipelines-with-reasoning-using-nvidia-llama-nemotron-models/
23. Enhancing Dialogue State Tracking Models through LLM-backed User-Agents Simulation - arXiv, accessed April 8, 2026, https://arxiv.org/html/2405.13037v1
24. LLMs vs. Traditional Search: What's Changing in SEO + How to Adapt - HawkSEM, accessed April 8, 2026, https://hawksem.com/blog/llms-vs-traditional-search-results/
25. Should You Use LLMs for Reranking? A Deep Dive into Pointwise, Listwise, and Cross-Encoders - ZeroEntropy, accessed April 8, 2026, https://www.zeroentropy.dev/articles/should-you-use-llms-for-reranking-a-deep-dive-into-pointwise-listwise-and-cross-encoders
26. How DoorDash leverages LLMs for better search retrieval, accessed April 8, 2026, https://careersatdoordash.com/blog/how-doordash-leverages-llms-for-better-search-retrieval/
27. Optimizing search relevance at Instacart using hybrid retrieval | by Vinesh Gudla, accessed April 8, 2026, https://tech.instacart.com/optimizing-search-relevance-at-instacart-using-hybrid-retrieval-88cb579b959c
28. What Are Embeddings in AI? How They Power Search and RAG [2026] - Atlan, accessed April 8, 2026, https://atlan.com/know/what-are-embeddings-ai-search/
29. Pinterest's CDC-Powered Ingestion Slashes Database Latency from 24 Hours to 15 Minutes, accessed April 8, 2026, https://www.infoq.com/news/2026/02/pinterest-cdc-db-ingestion/
30. A Survey on E-Commerce Learning to Rank - arXiv, accessed April 8, 2026, https://arxiv.org/html/2412.03581v1
31. How Good are LLM-based Rerankers? An Empirical Analysis of State-of-the-Art Reranking Models - arXiv, accessed April 8, 2026, https://arxiv.org/html/2508.16757v1
32. How Good are LLM-based Rerankers? An Empirical Analysis of State-of-the-Art

Reranking Models - ACL Anthology, accessed April 8, 2026, https://aclanthology.org/2025.findings-emnlp.305.pdf
33. ATTENTION IN LARGE LANGUAGE MODELS YIELDS EFFICIENT ZERO-SHOT RE-RANKERS - OpenReview, accessed April 8, 2026, https://openreview.net/pdf?id=yzloNYH3QN
34. FIRST: Faster Improved Listwise Reranking with Single Token Decoding - IBM Research, accessed April 8, 2026, https://research.ibm.com/publications/first-faster-improved-listwise-reranking-with-single-token-decoding
35. The Hidden Cost of Search Latency: Why Milliseconds Cost Millions - RETISIO.com, accessed April 8, 2026, https://www.retisio.com/hidden-cost-of-search-latency/
36. LLM-Driven Product Recommendations - Emergent Mind, accessed April 8, 2026, https://www.emergentmind.com/topics/llm-driven-product-recommendations
37. Durable multi-agentic AI architecture with Temporal, accessed April 8, 2026, https://temporal.io/blog/using-multi-agent-architectures-with-temporal
38. Designing Multi-Agent Intelligence - Microsoft for Developers, accessed April 8, 2026, https://developer.microsoft.com/blog/designing-multi-agent-intelligence
39. Language Agent Tree Search Unifies Reasoning, Acting, and Planning in Language Models - arXiv, accessed April 8, 2026, https://arxiv.org/html/2310.04406v3
40. Language Agent Tree Search Unifies Reasoning, Acting, and Planning in Language Models - arXiv, accessed April 8, 2026, https://arxiv.org/pdf/2310.04406
41. AI Agent Architecture Patterns: Single & Multi-Agent Systems - Redis, accessed April 8, 2026, https://redis.io/blog/ai-agent-architecture-patterns/
42. Customize agent workflows with advanced orchestration techniques using Strands Agents, accessed April 8, 2026, https://aws.amazon.com/blogs/machine-learning/customize-agent-workflows-with-advanced-orchestration-techniques-using-strands-agents/
43. Language Agent Tree Search Unifies Reasoning, Acting, and Planning in Language Models, accessed April 8, 2026, https://proceedings.mlr.press/v235/zhou24r.html
44. Building an AI-Powered Product Discovery Tool using Function Calling | by Rehmat Sayany, accessed April 8, 2026, https://rehmat-sayany.medium.com/building-an-ai-powered-product-discovery-tool-using-function-calling-f005aff4b9ed
45. UCP vs ACP: What Brands Need to Know About the Agentic Commerce Infrastructure Wars, accessed April 8, 2026, https://agenticplug.ai/blog/universal-commerce-protocol-vs-agentic-commerce-protocol
46. MCP vs UCP: Understanding their roles in AI-driven commerce - Inriver, accessed April 8, 2026, https://www.inriver.com/resources/mcp-vs-ucp-ai-commerce/
47. Function calling using LLMs - Martin Fowler, accessed April 8, 2026, https://martinfowler.com/articles/function-call-LLM.html
48. Understanding MCP, ACP and UCP in agentic commerce: A merchant's guide, accessed April 8, 2026, https://commercetools.com/blog/understanding-mcp-acp-ucp-in-agentic-comm

erce
49. A Strategic Analysis of Universal Commerce Protocol (UCP) and Agentic Commerce Protocol (ACP) - Advanced Web Ranking, accessed April 8, 2026, https://www.advancedwebranking.com/blog/ucp-acp-protocols-analysis
50. OpenAI ACP vs Google UCP: What's the Difference? - Cahoot.ai, accessed April 8, 2026, https://www.cahoot.ai/openai-acp-vs-google-ucp/
51. Universal Commerce Protocol (UCP) Explained - Salesforce, accessed April 8, 2026, https://www.salesforce.com/commerce/ai/agentic-commerce/universal-commerce-protocol/
52. Under the Hood: Universal Commerce Protocol (UCP) - Google Developers Blog, accessed April 8, 2026, https://developers.googleblog.com/under-the-hood-universal-commerce-protocol-ucp/
53. Universal Commerce Protocol - Universal Commerce Protocol (UCP), accessed April 8, 2026, http://ucp.dev/
54. Architecting for agentic commerce - Vaimo, accessed April 8, 2026, https://www.vaimo.com/blog/architecting-for-agentic-commerce/
55. Stop AI Agent Hallucinations: 4 Essential Techniques - DEV Community, accessed April 8, 2026, https://dev.to/aws/stop-ai-agent-hallucinations-4-essential-techniques-2i94
56. Advanced RAG Techniques for High-Performance LLM Applications - Neo4j, accessed April 8, 2026, https://neo4j.com/blog/genai/advanced-rag-techniques/
57. Minimize AI hallucinations and deliver up to 99% verification accuracy with Automated Reasoning checks: Now available | AWS News Blog, accessed April 8, 2026, https://aws.amazon.com/blogs/aws/minimize-ai-hallucinations-and-deliver-up-to-99-verification-accuracy-with-automated-reasoning-checks-now-available/
58. [2503.04830] Cite Before You Speak: Enhancing Context-Response Grounding in E-commerce Conversational LLM-Agents - arXiv, accessed April 8, 2026, https://arxiv.org/abs/2503.04830
59. Citations with Amazon Nova understanding models | Artificial Intelligence - AWS, accessed April 8, 2026, https://aws.amazon.com/blogs/machine-learning/citations-with-amazon-nova-understanding-models/
60. Cite Before You Speak: Enhancing Context-Response Grounding in E-commerce Conversational LLM-Agents - arXiv, accessed April 8, 2026, https://arxiv.org/html/2503.04830v3
61. [PDF] Cite Before You Speak: Enhancing Context-Response Grounding in E-commerce Conversational LLM-Agents | Semantic Scholar, accessed April 8, 2026, https://www.semanticscholar.org/paper/f2f1a200f179778319ac1cf8ea8e36fabfa1562b
62. Generating Comparative Table by LLM-Based Product Review ..., accessed April 8, 2026,

https://www.researchgate.net/publication/398287263_Generating_Comparative_Table_by_LLM-Based_Product_Review_Summarization
63. Why Product Comparison Works For Ecommerce Business? | SellersCommerce, accessed April 8, 2026, https://www.sellerscommerce.com/blog/ecommerce-product-comparison-guide/
64. Normalized Discounted Cumulative Gain (NDCG) explained - Evidently AI, accessed April 8, 2026, https://www.evidentlyai.com/ranking-metrics/ndcg-metric
65. Amazon-ESCI - Kaggle, accessed April 8, 2026, https://www.kaggle.com/datasets/marquis03/amazon-esci
66. Evaluating E-Commerce Search Engines with Amazon ESCI Dataset | r2decide Blog, accessed April 8, 2026, https://r2decide.com/blog/evaluating-e-commerce-search-engines
67. FACTS Grounding Leaderboard | Kaggle, accessed April 8, 2026, https://www.kaggle.com/benchmarks/google/facts-grounding
68. FACTS Grounding: A new benchmark for evaluating the factuality of large language models, accessed April 8, 2026, https://deepmind.google/blog/facts-grounding-a-new-benchmark-for-evaluating-the-factuality-of-large-language-models/
69. Enhancing Factual Accuracy and Citation Generation in LLMs via Multi-Stage Self-Verification - arXiv, accessed April 8, 2026, https://arxiv.org/html/2509.05741v1
70. Sub-100ms Discovery: Why Retrieval Speed is the Agent Bottleneck - Shaped.ai, accessed April 8, 2026, https://www.shaped.ai/blog/sub-100ms-discovery-why-retrieval-speed-is-the-agent-bottleneck
71. AI-driven personalization in consumer goods and retail: A technical analysis - World Journal of Advanced Research and Reviews, accessed April 8, 2026, https://wjarr.com/sites/default/files/fulltext_pdf/WJARR-2025-1639.pdf
72. The Cold Start Shopper Problem: Providing Personalization in a Private World - Total Retail, accessed April 8, 2026, https://www.mytotalretail.com/article/the-cold-start-shopper-problem-providing-personalization-in-a-private-world/
73. Towards Personalized Conversational Sales Agents: Contextual User Profiling for Strategic Action - arXiv, accessed April 8, 2026, https://arxiv.org/html/2504.08754v4
74. Enhancing Generative AI Shopping Assistants through Advanced Multi-Attribute Decision Making Technique - ResearchGate, accessed April 8, 2026, https://www.researchgate.net/publication/395348639_Enhancing_Generative_AI_Shopping_Assistants_through_Advanced_Multi-Attribute_Decision_Making_Technique
75. Deep Dive into Advanced Preference Elicitation Techniques, accessed April 8, 2026, https://sparkco.ai/blog/deep-dive-into-advanced-preference-elicitation-techniques

76. PSCon: Product Search Through Conversations - arXiv, accessed April 8, 2026, https://arxiv.org/html/2502.13881v3
77. PSCon: Toward Conversational Product Search - arXiv, accessed April 8, 2026, https://arxiv.org/html/2502.13881v1